



PostgreSQL Tuning Basics

Warsaw PostgreSQL Meetup / 23.11.2019

Tomáš Vondra

tomas.vondra@2ndquadrant.com / tomas@pgaddict.com

© 2019 Tomas Vondra, under Creative Commons Attribution-ShareAlike 3.0

<http://creativecommons.org/licenses/by-sa/3.0/>

Agenda

1) basic configuration

- shared_buffers
- (maintenance_)work_mem
- max_connections
- effective_cache_size

2) checkpoint tuning

- checkpoint_segments (timeout / completion_target)
- max_wal_size
- bgwriter (delay / ...)

3) autovacuum tuning

- scale factor, limit, ...

4) other config options

- wal_level
- synchronous_commit
- default_statistics_target
- effective_io_concurrency

5) a little bit about hardware / OS

- ... the whole time

Sources

PostgreSQL 9.0 High Performance (Gregory Smith)

- exhaustive analysis of the topic
- more or less basis for this workshop

PostgreSQL 9 High Availability (Shaun M. Tomas)

- not really about tuning, but HA is “related topic”
- hardware planning, performance triage, ...

What Every Programmer Should Know About Memory



- Ulrich Drepper, Red Hat
- <http://www.akkadia.org/drepper/cpumemory.pdf>
- low-level features of CPU and RAM

Righting Your Writes (Greg Smith)

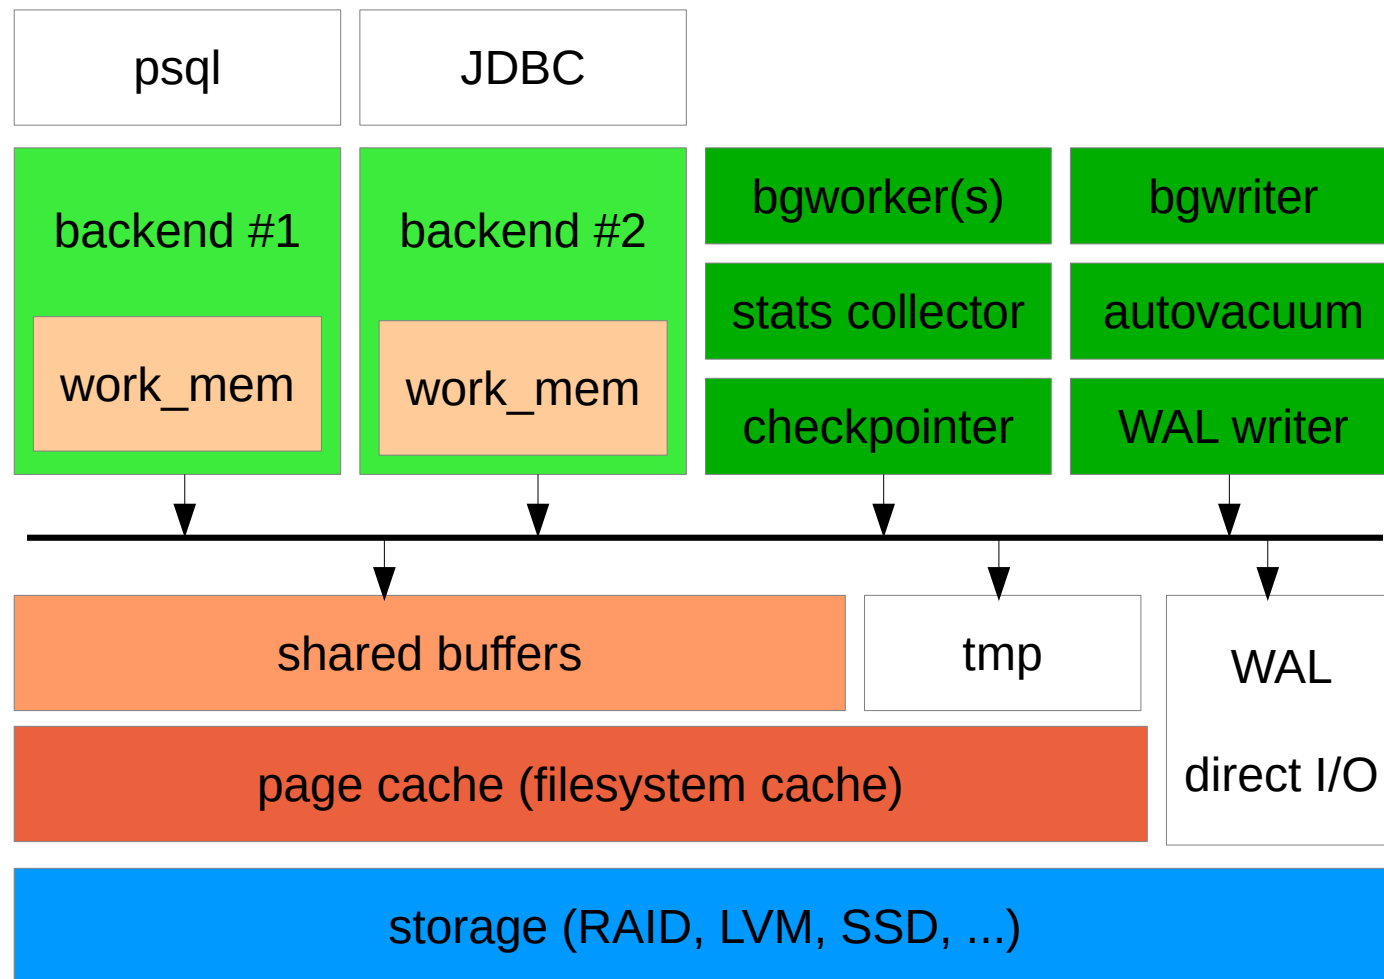
- <http://2ndquadrant.com/media/pdfs/talks/RightingWrites.pdf>

PostgreSQL Wiki

- https://wiki.postgresql.org/wiki/Tuning_Your_PostgreSQL_Server



PostgreSQL architecture



basic configuration

shared_buffers

- paměť vyhrazená pro databázi
- prostor sdílený všemi databázovými procesy
- cache “bloků” z datových souborů (8kB)
 - částečně duplikuje page cache (double buffering)
- bloky se dostávají do cache když ...
 - backend potřebuje data (SQL dotaz, autovacuum, ...)
- bloky se dostávají z cache když
 - nedostatek místa v cache (LRU)
 - průběžně (background writer)
 - checkpoint
- bloky mohou být čisté nebo změněné (“dirty”)

shared_buffers

- default 128MB (used to be 32MB before 9.3)
 - goal “has to start everywhere” - low default values
 - 32MB limit is motivated by kernel limits (SHMALL)
 - 9.3 allocates shmem differently, SHMALL irrelevant
 - 128MB better, but still conservative (small systems)
 - can benefit from (explicit) huge pages
- What is optimal value?
 - high “cache hit ratio”, without wasting memory
 - larger sizes -> higher overhead, double buffering
- Why DB won't check available RAM and pick “optimal” size?
 - depends on workload (how app uses the DB) and “active set”
 - might be sharing RAM with other stuff (appserver, ...)

shared_buffers

- iterative monitoring-based tuning
 1. pick conservative initial value (1GB?)
 2. measure important metric
 - cache hit ratio (viz. `pg_stat_bgwriter`)
 - usage of shared buffers (`pg_buffercache`)
 - eviction of dirty buffers (`pg_stat_bgwriter`, checkpoints)
 - latence of operations (queries), maintenance, data loads, ...
 3. increase `shared_buffers` size (e.g. 2x)
 4. measure important metrics again
 - Did they improve? Repeat the shared buffer size increase (2x)
 - reduce the size again and finish
- reproducible application benchmark (not a stress test)
 - the same thing, but you can iterate much faster

shared_buffers

- pg_buffercache
 - <http://www.postgresql.org/docs/devel/static/pgbuffercache.html>
 - extension available with PostgreSQL (usually in -contrib package)
 - adds a new system view (list of blocks in shared buffer cache)

```
CREATE EXTENSION pg_buffercache;
```

```
SELECT
```

```
    datname,
```

```
    usagecount,
```

```
    COUNT(*) AS buffers,
```

```
    COUNT(CASE WHEN isdirty THEN 1 ELSE NULL END) AS dirty
```

```
FROM pg_buffercache JOIN pg_database d
```

```
    ON (relidatabase = d.oid)
```

```
GROUP BY 1, 2
```

```
ORDER BY 1, 2;
```

work_mem

- memory limit for operations (sorts, hash tables, ...)
 - default 4MB (very conservative, fine for OLTP)
 - one query can do multiple operations -> multiple buffers
 - affects planning (query costing, possibility of plans)
 - some operations don't fully respect (Hash Aggregate)
- when exceeded, a temporary file is used
 - not necessarily a slowdown (can stay in page cache)
 - may use a different algorithm (quick-sort, merge sort)
- optimální value depends on
 - RAM available (after subtracting shared buffers)
 - number of parallel queries
 - query complexity (OLTP vs. OLAP/BI)

work_mem

- example
 - system has (RAM – shared_buffers) available memory
 - we don't want to use everything (page cache, OOM atd.)
 - expect all connections are active (consider connection pool)
 - expect each query uses 2 * work_mem

`work_mem = 0.25 * (RAM - shared_buffers) / max_connections / 2;`

- “spojené nádoby FIXME” (fewer queries -> higher work_mem possible)
- alternative approach
 - look at slow queries
 - would they benefit from increasing work_mem / how much?
 - check if that risks OOM and then change config

work_mem

- work_mem not necessarily the same for everyone
- can be modified per session

```
SET work_mem = '1TB';
```

- can be modified per user

```
ALTER USER webuser SET work_mem = '8MB';  
ALTER USER dwhuser SET work_mem = '128MB';
```

- can be modified per database

```
ALTER DATABASE webapp SET work_mem = '8MB';  
ALTER DATABASE dwh SET work_mem = '128MB';
```

- <http://www.postgresql.org/docs/devel/static/sql-alteruser.html>
- <http://www.postgresql.org/docs/devel/static/sql-alterdatabase.html>

maintenance_work_mem

- similar to work_mem, but for “maintenance” operations
 - CREATE INDEX, REINDEX, VACUUM, REFRESH
- default 64MB – not bad, but increase can be quite beneficial
 - e.g. REINDEX of large tables etc.
- may have significant impact, but not necessarily “the more is better”

```
test=# set maintenance_work_mem = '4MB';
test=# create index test_1_idx on test(i);
CREATE INDEX
Time: 27076,920 ms
```

```
test=# set maintenance_work_mem = '64MB';
test=# create index test_1_idx on test(i);
CREATE INDEX
Time: 39468,621 ms
```

max_connections

- default value 100 is a bit too high in many cases
 - assumes many connections are inactive
 - may not be true, backends will interfere and cause slowdown
 - context switches, lock contention, disk contention, more RAM used, cache line contention (CPU caches), ...
 - results in lower performance / throughput, latencies, ...
- rough “traditional” formula
$$((\text{core_count} * 2) + \text{effective_spindle_count})$$
- better to use lower value and a connection pool (e.g. pgbouncer)

https://wiki.postgresql.org/wiki/Number_Of_Database_Connections

wal_level

- determines which information need to be written to Write Ahead Log
- multiple levels, adding more and more information
- `minimal`
 - local recovery only (crash, immediate shutdown)
 - may skip WAL for some commands (CREATE TABLE AS, CREATE INDEX, CLUSTER, COPY into a table created in the same transaction)
- `replica` (10+)
 - WAL archiving (log-file shipping replication, `warm_standby`)
 - read-only standby
- `logical`
 - allows logical replication (interprets WAL log)

effective_cache_size

- default 4GB, but does not directly allocate anything
- simply a “hint” for the query planner
 - How likely is it that block “X” is in memory / no disk read needed?
 - What fraction of blocks will I need to read from the disk?
- good formula
$$(\text{shared buffers} + \text{page cache}) * \text{fraction}$$
- page cache is an estimate
 - remaining RAM without kernel memory, work_mem, other apps ...
- often used fractions
 - 0.75 - aggressive (a lot of sharing of data between backends)
 - $1/\text{max_connections}$ – defensive, backends on different subsets of data
- usually not worth spending a lot of time on
 - reasonable default, increasing has small impact (compared to other options)

checkpoint tuning

<https://blog.2ndquadrant.com/basics-of-tuning-checkpoints/>

CHECKPOINT

- WAL
 - split into 16MB segments (by default)
 - limited number of segments, recycling
- COMMIT
 - write into a transaction log (WAL) + fsync
 - sequential writes (efficient on most hardware)
 - modify data files in shared_buffers (no immediate disk write)
- CHECKPOINT
 - after “filling” WAL or timeout (checkpoint_timeout)
 - writes out changes from shared buffers to data files
 - write to page cache + fsync at the end
 - checkpoint_flush_after helps to remove “spikes”

CHECKPOINT

- checkpoints need to be done with “proper frequency”
 - too often – prevents optimizations (merging writes, ordering writes)
 - too rarely – long recovery, have to keep more WAL segments
- two basic “triggers” for checkpoint
 - expiration of a time limit (checkpoint_timeout)
 - generating too much WAL (checkpoint_segments / max_wal_size)

$(3 * \text{checkpoint_segments}) \sim \text{max_wal_size}$

checkpoint_timeout

- checkpoint_timeout
 - maximum distance between checkpoints
 - default 5 minut (fairly aggressive), maximum 1 day
- rough upper limit on recovery time (but not quite)
 - recovery is often faster (just writes to data files)
 - but not necessarily
 - recovery is single-threaded
 - may not have anything in memory (reboot)

checkpoint_completion_target

- up to 8.2 problem with I/O spiked during checkpoint
 - write everything at once + fsync
 - goal:
 - spread writes to page cache in time
 - finish writes with enough remaining time for kernel to do flushes in the background (final fsync fast)
 - works both with “timed” and “xlog” checkpoints
 - checkpoint_flush_after – alternative solution

checkpoint tuning

- pg_stat_bgwriter

```
SELECT checkpoints_timed, checkpoints_req  
FROM pg_stat_bgwriter;
```

checkpoints_timed		checkpoints_req
-----+-----		
201		159

- vast majority of checkpoints should be “timed”
- goal is to minimize checkpoints_req value
 - can’t be 100% (shutdown, CREATE DATABASE, ...)

bgwriter

- background writer (bgwriter)
 - background process regularly walking shared buffers, evicting unused ones
 - makes sure there are enough clean (not modified) buffers for queries
- pg_stat_bgwriter
 - system catalog (global) with bgwriter statistics
 - number of blocks written for various reasons (and other metrics)
 - buffers_alloc – blocks loaded into shared buffers
 - buffers_checkpoint – written out by checkpointer
 - buffers_clean – written out by bgwriter
 - buffers_backend – written out by backends (impact on queries)

```
SELECT
    now(),
    buffer_checkpoint, buffer_clean, buffer_backend, buffer_alloc
FROM pg_stat_bgwriter;
```

bgwriter (delay / ...)

- alternative approach to sizing shared buffers
 - smaller shared buffers + more aggressive background eviction
 - often you can't have sufficiently large shared buffers
- `bgwriter`
 - monitor number of buffers needed by backends per interval
 - evict a multiple of the number (in the background)
- `bgwriter_delay = 200ms`
 - delay between runs of bgwriter process
- `bgwriter_lru_multiplier = 2.0`
 - multiple of pages needed in previous round
- `bgwriter_lru_maxpages = 100`
 - max number of pages written out in each round


autovacuum tuning

[**https://blog.2ndquadrant.com/autovacuum-tuning-basics/**](https://blog.2ndquadrant.com/autovacuum-tuning-basics/)

autovacuum options

- `autovacuum = on`
- `autovacuum_work_mem = -1` (`maintenance_work_mem`)
- `autovacuum_max_workers = 3`
- `autovacuum_naptime = 1min`
- **`autovacuum_vacuum_threshold = 50`**
- **`autovacuum_analyze_threshold = 50`**
- **`autovacuum_vacuum_scale_factor = 0.2`**
- **`autovacuum_analyze_scale_factor = 0.1`**
- ...
- `autovacuum_vacuum_cost_delay = 20ms`
- **`autovacuum_vacuum_cost_limit = -1`** (`vacuum_cost_limit = 200`)
- `vacuum_cost_page_hit = 1`
- `vacuum_cost_page_miss = 10`
- `vacuum_cost_page_dirty = 20`

autovacuum = off



http://en.wikipedia.org/wiki/Nuclear_explosion

launcher and workers

- `autovacuum_naptime = 1min`
 - interval between runs of “autovacuum launcher” process
 - in each interval, the launcher will try to start a worker on each db
 - interval between starting autovacuum workers ($1\text{min} / \text{num of DBs}$)
 - else – interval between autovacuum runs on a given DB
- `autovacuum_workers = 3`
 - number of “worker” processes doing the actual work is limited
 - if all workers are busy, new can't be started

autovacuum thresholds

- how does the autovacuum system which tables need cleanup?
 - monitoring of number of deleted/modified rows
 - after exceeding a limit, the table is considered for cleanup

`modified_rows > (threshold + total_rows * scale_factor)`

- `autovacuum_vacuum_threshold = 50`
- `autovacuum_vacuum_scale_factor = 0.2`
- `autovacuum_analyze_threshold = 50`
- `autovacuum_analyze_scale_factor = 0.1`
- row counts come from system catalogs (`pg_class`, `pg_stat_all_tables`)

autovacuum limit

- `autovacuum_vacuum_cost_limit`
 - globální limit, sdílený všemi autovacuum worker procesy
 - zvýšení `autovacuum_max_workers` většinou nic neřeší (je jich víc ale pracují pomaleji)
- lze předefinovat pro jednotlivé tabulky

```
ALTER TABLE t SET (autovacuum_vacuum_cost_limit = 1000);
```

- tabulka (resp. autovacuum worker) je vyjmuta z globálního limitu a limit je aplikován na samostatného workera
- ale stále to nezaručuje že volný worker bude k dispozici

autovacuum throttling

- autovacuum workers should not use too much resources (I/O, CPU)
- throttling activity over time, based on basic operations
 - `vacuum_cost_page_hit = 1` # read from cache
 - `vacuum_cost_page_miss = 10` # read from OS
 - `vacuum_cost_page_dirty = 20` # modified
- time is divided into small intervals, with budget per interval
 - `autovacuum_vacuum_cost_delay = 20ms`
 - `autovacuum_vacuum_cost_limit = -1 (200)`
- so there's total “per second” budget 10000 (= 50 x 200)
 - 10000 cache hits / second
 - 1000 reads / second => 8MB/s
 - 500 writes / second => 4MB/s

autovacuum tuning

- DON'T DISABLE AUTOVACUUM!
 - Seriously. Don't repeat our mistakes.
- increase the throttling limits (4/8 MB/s is way too low)
 - increase cost_limit, decrease cost_delay
 - depends on available hardware resources (CPU, I/O)
 - current systems should handle 10x that
- trigger autovacuum more often
 - “If it hurts, do it more often.”
 - 10% is fine on 10GB table, not so much on 1TB one
 - Don't overdo it, a bit of bloat is natural / beneficial.
- maybe increase number of workers

autovacuum fails

- Triggering autovacuum more often can be making things worse.
- If there's a long transaction (forgotten session, prepared transaction, ...)
 - autovacuum can't actually cleanup anything
 - it'll be triggered over and over (CPU utilization, I/O traffic, ..)
- autovacuum can also do ANALYZE to collect stats
 - both phases are throttled
 - VACUUM cancels itself when there's lock request on the table
 - ANALYZE can't cancel itself – may block DDL

logging and monitoring

logging & monitoring

- important options
 - `log_line_prefix` (string)
 - `log_min_duration_statement` (integer)
 - `log_checkpoints` (boolean)
 - `log_temp_files` (integer)
 - `log_lock_waits` (integer)
 - `log_auto_vacuum_min_duration` (integer)
- interesting tools
 - <http://dalibo.github.io/pgbadger/>

auto_explain

- `auto_explain.log_min_duration` (integer)
- `auto_explain.log_analyze` (boolean)
- `auto_explain.log_buffers` (boolean)
- `auto_explain.log_timing` (boolean)
- `auto_explain.log_triggers` (boolean)
- `auto_explain.log_verbose` (boolean)
- `auto_explain.log_format` (enum)
- ... další volby ...

<https://www.postgresql.org/docs/current/auto-explain.html>

pg_stat_statements

- userid
- dbid
- queryid
- query
- calls
- total_time
- rows
- shared_blks_hit
- shared_blks_read
- shared_blks_dirtied
- shared_blks_written
- local_blks_hit
- local_blks_read
- local_blks_dirtied
- local_blks_written
- temp_blks_read
- temp_blks_written
- blk_read_time
- blk_write_time

other config options

Durability tuning

safe

- synchronous_commit = off
- checkpoint_segments / max_wal_size = high number
- unlogged tables (lost after DB crash, not replicated)

unsafe

- fsync = off
- full_page_writes = off
- unlogged tables (lost after DB crash, not replicated)

<https://www.postgresql.org/docs/current/non-durability.html>

synchronous_commit

- should we wait for confirmation tuning?
 - “durability tuning” long before the NoSQL hype
 - still fully transactional / ACID
- up to 9.0 only on / off options
- 9.1 added sync replication – many more options
 - on (default) – wait for commit confirmation
 - off – don’t wait for local WAL confirmation
 - local – do not wait for a replica, local WAL is enough
 - remote_write – wait for write into WAL on a replica
 - remote_apply – applied on replica (visible)
- can be set “per transaction”
 - important transactions “on”, less important “local”

<https://www.postgresql.org/docs/current/runtime-config-wal.html>

wal_log_hints

- MVCC
 - each row has ID of two transactions – INSERT / DELETE
 - when reading data, we need to check visibility of those xids
 - expensive (CPU), hint bits are “flags” caching the results
 - not necessarily WAL-logged (can be recalculated)
 - problem after recovery on replicas (after failover / hot standby)
 - hint bits not set, everything has to be checked from scratch
 - data checksums enable this automatically
- odkazy
 - http://en.wikipedia.org/wiki/Multiversion_concurrency_control
 - <http://www.postgresql.org/docs/current/static/mvcc-intro.html>
 - <http://momjian.us/main/writings/pgsql/internalpics.pdf>
 - <http://momjian.us/main/writings/pgsql/mvcc.pdf>

random_page_cost

- our optimization is based on computing “cost” of queries
 - amount of resources (CPU, I/O consumed by the execution)
 - lower cost => can execute faster => lower duration
- five basic cost parameters determining cost of basic operations
 - seq_page_cost = 1
 - random_page_cost = 4
 - cpu_tuple_cost = 0.01
 - cpu_index_tuple_cost = 0.005
 - cpu_operator_cost = 0.0025
- difficult to verify changes to those values
 - may improve one query, hurt other
- the most common thing is tweaking random_page_cost
 - on SSD, big RAID arrays maybe lower to 2, or even 1.5

statement_timeout

- from time to time you'll get "runaway query"
 - e.g. cartesian product generating 100 trillions of rows
 - using a lot of CPU or I/O (or both)
 - affects other activity on the system (user queries, vacuuming)
- it's better kill / fix such queries, because they'll not finish anyway
- statement_timeout
 - limit on maximal query duration (milliseconds)
 - affects "everything" (data loads, ...)
 - just like work_mem etc. can be set per user / db
- alternative
 - cron skript (allows e.g. matching queries by regexp, ...)

temp_file_limit

- another way to limit “greedy” queries
 - if query requires too much temporary files
- usually means the query will run for a long time anyway
 - so statement_timeout will kill it eventually
 - but until then it'll put pressure on the I/O subsystem
 - may push all interesting data from page cache
 - may interact with kernel write cache config (dirty_bytes, background_dirty_bytes)